

Overall verdict

This is a major improvement over the rejected version and it does implement the strategy we discussed. The new paper is now much more clearly about **systemic financial stress, volatility connectedness, market fragility, and market-implied regimes**, rather than about the construction and validation of the geopolitical shock index. That substantially reduces the overlap with the BIR paper.

However, **I would not submit this version to Global Finance Journal yet.** It is directionally promising, but it currently has several issues that a careful GFJ reviewer could use to recommend rejection or a very difficult revision. The most serious problems are not the idea; the idea is strong. The main problems are **data consistency, internal contradictions, identification fragility, HMM/TCI implementation credibility, incomplete references, and overclaiming of statistically weak results.**

My reviewer-style recommendation would be:

Promising paper, potentially suitable for GFJ after substantial revision, but not submission-ready in the current form.

1. Did the paper do what we discussed?

Yes, largely. It follows the exact architecture we discussed:

Previously discussed component	Implemented in the paper?	Assessment
Use geopolitical shocks only as an external trigger	Yes	Good strategic distancing from BIR
Build market-based systemic fragility measures	Yes	Strong move
Add volatility connectedness / spillover channel	Yes, via TCI	Good, but TCI implementation needs strengthening
Build a European Market Fragility Index	Yes, EMFI via PCA	Good, but requires better validation and out-of-sample treatment
Add market-implied stress detection	Yes, HMM	Very useful, but currently too mechanically tied to EMFI components
Ask “when do geopolitical shocks become systemic?”	Yes	Excellent new framing
Add state dependence	Yes	This should become the paper’s main contribution

Previously discussed component	Implemented in the paper?	Assessment
Add absorbed/localized/systemic taxonomy	Yes	Interesting and publishable if made internally consistent
Move away from BIR measurement contribution	Mostly	Better, but still some shock-construction inconsistencies

The abstract already states the new research question clearly: the paper studies when geopolitical shocks translate into systemic financial stress using 19 European equity markets, EMFI, HMM stress regimes, and local projections. That is exactly the right direction.

The introduction also now uses the “three-layer” structure we discussed: geopolitical shock trigger, market-based fragility measures, and an HMM regime classifier independent of geopolitical event dates. This is a much better conceptual contribution than the previous “geopolitical shocks affect volatility” paper.

So, conceptually, the paper is now on the right path.

2. Fit with Global Finance Journal

The fit is **much better than before**.

The previous paper was vulnerable because it looked like a methodological shock-measurement paper with a volatility application. This new version is closer to **global finance**, because it studies:

- European equity-market integration;
- systemic fragility;
- volatility connectedness;
- cross-market tail co-exceedance;
- market-implied stress regimes;
- state-dependent amplification of geopolitical shocks.

That is much closer to what a global/international finance audience would expect.

For GFJ, the best positioning is not:

“We develop a new geopolitical shock measure.”

It should be:

“We show that geopolitical shocks become financially systemic only when they interact with pre-existing market fragility and volatility connectedness.”

That is the paper’s strongest claim.

3. The paper's main strengths

3.1 The new research question is much stronger

The central question — “**when do geopolitical shocks become systemic financial stress events?**” — is excellent. It is sharper, more policy-relevant, and much more distinct from BIR than the old question. The introduction clearly distinguishes a shock that affects one market from a shock that pushes many European markets into tail states and raises connectedness.

That distinction is valuable. It shifts the paper from a standard event-response exercise to a **systemic risk transmission** paper.

3.2 The EMFI idea is promising

The paper constructs EMFI from volatility stress, tail-volatility breadth, average correlation, and TCI. The first principal component explains 58.3% of joint variance, with balanced positive loadings across all four components.

This is a sensible composite. It captures both acute volatility stress and cross-market integration. The discussion of the two-cluster covariance structure is also honest and useful: volatility/tail breadth form one block, while average correlation/TCI form another.

That is much better than pretending the first PC is a perfect one-dimensional systemic-risk factor.

3.3 The HMM layer is a good addition

The HMM is a useful way to introduce market-implied stress regimes without using geopolitical shock dates. The paper explicitly states that the HMM is estimated only from market-based fragility measures and contains no information about geopolitical shocks.

That is exactly the logic we wanted: the market-implied stress detector should be independent of the geopolitical shock trigger.

3.4 State dependence is the most interesting finding

The strongest empirical idea is that geopolitical shocks are mostly absorbed in calm markets but amplified in fragile markets. The smooth-transition specification is conceptually good, and the reported amplification magnitudes are economically meaningful: EMFI amplification of 6.4x, TCI amplification of 10.3x, tail-breadth amplification of 4.1x, and Pstress amplification of 7.8x.

This should become the paper's core contribution.

3.5 The event taxonomy is attractive

The absorbed/localized/systemic taxonomy is potentially very publishable. The paper reports that 50.6% of shocks are absorbed, 22.1% are localized, and 27.3% are systemic.

That is a more nuanced and interesting message than “geopolitical shocks increase volatility.” It says:

Most geopolitical shocks do not become systemic; the ones that do are those interacting with fragile market states.

That is the right paper.

4. Major concerns that must be fixed before submission

4.1 The country sample is inconsistent with the previous/BIR shock series

This is one of the most serious problems.

The current paper says the shock series is inherited from the companion BIR paper, but the country list is different. The current paper includes **Czech Republic, Denmark, and Switzerland**, while the earlier 19-country BIR/rejected-paper setup included countries such as **Bulgaria, Croatia, and the United Kingdom**.

Table 1 in the current paper lists Austria, Belgium, Czech Republic, Denmark, Finland, France, Germany, Greece, Hungary, Ireland, Italy, Netherlands, Norway, Poland, Portugal, Romania, Spain, Sweden, and Switzerland.

This creates a major credibility problem. If the shock series is “inherited” from the companion paper, the exact country set must match. If you changed the equity-market sample, then the shock series must be regenerated for the new 19 countries.

A reviewer will ask:

Are the shocks available for the same 19 countries used in the equity-market panel?

If not, what exactly is MaxShock measuring?

Recommendation: Either restore the original 19-country sample used in the BIR paper, or regenerate the shock series for the new 19-country sample. Do not say the country assignment is inherited from BIR unless the countries truly match.

4.2 The shock-construction description conflicts with the BIR paper

The current paper describes the conflict intensity index as being computed from “CAMEO action codes weighted by the Goldstein scale.”

But the BIR paper, as previously discussed, was based on **GDEL Document 2.0 conflict coverage**, not a CAMEO/Goldstein event-score index. This is not a small wording issue. It changes the object being measured.

You also say the FDR procedure is applied cross-sectionally and “ensures that only days with statistically significant simultaneous exceedances across countries are retained.” That is stronger than the more careful interpretation we previously developed, where q-values provide multiple-testing transparency and continuous intensity, while the binary event rule uses raw tail probabilities.

Recommendation: In this paper, do not describe the shock-construction mechanics in detail. Use a short summary and refer to the companion paper. Something like:

“We use the validated conflict-shock series from the companion measurement paper. The series is based on country-level conflict-related news coverage from GDELT Document 2.0, transformed into tail-based shock intensities using EVT and multiple-testing adjustments. We use the resulting country-day shock intensities as external triggers.”

Then add a footnote:

“The present paper does not contribute to shock construction; it takes the validated shock series as given.”

That protects this paper from BIR overlap and avoids methodological inconsistency.

4.3 COVID-19 should not be treated as a geopolitical conflict shock

The introduction opens with “the COVID-19 shock of February–March 2020” as a geopolitical event. This is risky.

COVID-19 is an enormous systemic financial stress event, but it is not naturally a **geopolitical conflict shock**. If your geopolitical shock series flags COVID-related days, that is itself a problem: it may mean the shock series is detecting general crisis news rather than geopolitical conflict.

This issue could severely damage the paper’s credibility.

Recommendation: Treat COVID as a **market-stress benchmark**, not as a geopolitical shock. Use it to validate EMFI/HMM, but do not include it as a geopolitical treatment event.

A better structure:

- COVID-19: used to show that EMFI/HMM detect systemic market stress.
- Ukraine invasion: primary geopolitical systemic event.
- Hamas-Israel / Iran-Israel / terrorist events: geopolitical shocks with varied market effects.
- COVID exclusion: not merely robustness, but baseline credibility.

If the main amplification result depends heavily on COVID shock days, then the paper is not cleanly about geopolitical shocks.

4.4 There is a direct contradiction about the Hamas-Israel event

The abstract says the 2023 Hamas-Israel shock is “absorbed without regime transition.”

But Section 7.2 says that the Hamas attack maps to 9 October 2023, has stress probability 0.296 on the shock day, then peaks at 1.0 on the following trading day, placing it in the **Systemic** category.

This is a serious internal inconsistency. A reviewer will catch it immediately.

Recommendation: Decide which result is correct and revise everywhere. If the post-window classification makes Hamas-Israel systemic, then the abstract and introduction must not call it absorbed.

If contemporaneous-only classification makes it absorbed, then the taxonomy should not use the full post-event window. You cannot have both.

4.5 The unconditional LP results are weak, but the conclusion overstates them

Section 6 says clearly that the contemporaneous EMFI response is **0.458 with $p = 0.174$** , and that none of the contemporaneous estimates reaches 5% significance.

But the conclusion says:

“The baseline panel local projection yields a positive and statistically significant impact coefficient on EMFI at $k = 0$.”

That is wrong relative to the results section.

This must be corrected. The abstract is actually more honest than the conclusion: it says baseline responses are positive but statistically imprecise.

Recommendation: Make the paper’s empirical hierarchy explicit:

1. Unconditional average effects are positive but weak/imprecise.
2. The main result is state-dependent amplification.
3. The taxonomy explains why unconditional effects are weak.

Do not call the unconditional EMFI coefficient statistically significant unless it is.

4.6 The TCI estimation is likely underpowered or overparameterized

This is a major econometric concern.

The paper estimates a rolling VAR on a 19-dimensional vector of absolute returns using a 100-day rolling window. A 19-variable VAR with even one lag has many parameters relative to only 100 observations per window. With $p = 1$, each equation has 19 lag coefficients plus intercept; across 19 equations, that is a very large parameter set for a short window. With $p = 2$ or $p = 3$, the problem becomes worse.

A GFJ reviewer familiar with connectedness methods may challenge this immediately.

Recommendation: Strengthen the TCI construction. Options:

1. Use a longer rolling window, e.g. 250 or 500 trading days.
2. Use a penalized VAR / LASSO-VAR connectedness estimator.
3. Use TVP-VAR connectedness.
4. Use a lower-dimensional factor-VAR on volatility factors.
5. Report sensitivity to window length and FEVD horizon.

6. Report whether the connectedness results are stable under a simpler correlation-network measure.

The current robustness with TCI windows of 60, 100, and 150 is useful, but it does not solve the overparameterization concern. A 60-day window is even more problematic.

4.7 The HMM may be detecting spikes, not regimes

The HMM systemic-stress state has a diagonal transition probability of 0.371 and an implied mean duration of about 1.6 trading days. That is very short for something called a “regime.”

This does not invalidate the approach, but the terminology should be adjusted. The HMM seems to detect **systemic stress episodes or spikes**, not persistent regimes.

Also, the Viterbi path identifies 143 distinct systemic-stress episodes with median duration one day. That reinforces the idea that these are spike classifications.

Recommendation: Replace “systemic stress regime” with one of:

- “market-implied systemic stress state”;
- “stress-state probability”;
- “latent stress episode”;
- “systemic stress classification.”

Also add robustness with:

- two-state HMM;
- Student-t HMM or robust emissions;
- diagonal vs full covariance;
- HMM estimated excluding COVID;
- HMM estimated only through pre-2020 and applied out-of-sample;
- filtered probabilities rather than smoothed probabilities.

This last point matters. If $P_{\text{stress},t}$ is based on full-sample smoothed probabilities, then it uses future information. For a monitoring/policy claim, you need **filtered** probabilities based only on information up to day t .

4.8 EMFI is ex post, but the paper sometimes interprets it as monitorable in real time

The paper says the EMFI components are standardized using full-sample means and standard deviations, explicitly because EMFI is an ex-post outcome variable, not a real-time detector.

That is fine. But later, the conclusion says monitoring EMFI and Pstress may be more informative for crisis prediction than monitoring geopolitical news flow.

That claim is too strong unless you construct EMFI and Pstress recursively or out-of-sample.

Recommendation: Either:

1. tone down the policy claim to “ex-post stress assessment,” or
2. build a real-time version using rolling/expanding standardization, rolling PCA loadings, and filtered HMM probabilities.

For GFJ, a real-time fragility indicator would be much more attractive.

4.9 The event taxonomy is promising but currently too endogenous

The taxonomy classifies events based on post-event EMFI and post-event Pstress. This is useful descriptively, but it is not an independent validation of the effect. It is partly a restatement of the outcome.

That is not fatal, but the paper should be explicit:

“The taxonomy is descriptive and outcome-based; it is not an ex ante classifier.”

If you want it to become stronger, estimate a predictive model:

$\Pr(\text{Systemic event}) = f(\text{pre-shock EMFI, pre-shock TCI, shock intensity, shock breadth})$

Then test whether pre-existing fragility predicts which geopolitical shocks become systemic. That would make the taxonomy far more valuable.

4.10 The references are not submission-ready

The manuscript contains many unresolved citations: [? ? ?], [?], and missing references throughout the literature review and methodology sections.

This is an automatic desk-rejection risk. It signals that the paper is still a draft.

Recommendation: Fix references before any external circulation. At minimum, the paper needs proper citations for:

- Caldara and Iacoviello geopolitical risk;
- Diebold-Yilmaz connectedness;
- generalized FEVD;
- Jordà local projections;
- Newey-West HAC;

- Hamilton Markov switching;
- Adrian-Brunnermeier CoVaR;
- state-dependent LPs;
- financial contagion/integration;
- recent geopolitical risk and stock-market connectedness papers.

5. The biggest strategic decision: what should be the main result?

Right now, the paper is trying to make three empirical claims:

1. geopolitical shocks raise EMFI/TCI on average;
2. shocks have larger effects in fragile states;
3. shocks can be classified as absorbed/localized/systemic.

The first claim is weak. The second and third are much more interesting.

For GFJ, I would make the paper about this:

Geopolitical shocks are not systematically destabilizing on average; they become systemic only when they interact with pre-existing market fragility.

This is actually a stronger and more original result than trying to force unconditional significance.

The paper should openly say:

“The average unconditional response is modest and imprecise. This is not a failure of the design; it is the central empirical fact. Most geopolitical shocks are absorbed. The important mechanism is state-dependent amplification.”

That is coherent and publishable.

6. Recommended revised contribution statement

I would rewrite the contribution section around three claims:

Contribution 1: A market-based systemic fragility framework

You construct EMFI and HMM stress probabilities from equity-market data, independent of geopolitical shock dates.

Contribution 2: State-dependent geopolitical risk transmission

You show that geopolitical shocks have limited average effects but large effects when markets are already fragile.

Contribution 3: A taxonomy of financially material geopolitical shocks

You classify shocks into absorbed, localized, and systemic events, showing that only a minority become systemic.

This is far better than saying “geopolitical shocks increase market fragility.”

7. What I would change in the abstract

The current abstract is mostly good, but it contains the Hamas contradiction and perhaps undersells the most interesting contribution.

A stronger abstract would be:

This paper asks when geopolitical shocks become systemic financial stress events. Using daily equity returns for 19 European markets from 2017 to 2025, we combine a validated high-frequency geopolitical conflict shock series with market-based measures of systemic fragility, including volatility stress, tail co-exceedance breadth, volatility connectedness, a PCA-based European Market Fragility Index, and an HMM-implied stress-state probability. Unconditional local projections show positive but statistically imprecise average responses, indicating that most geopolitical shocks are absorbed by markets. The main finding is state dependence: shocks that occur when market fragility is already elevated generate substantially larger responses in EMFI, volatility connectedness, and tail co-exceedance than shocks that occur in calm markets. An event taxonomy confirms that only a minority of geopolitical shocks become systemic stress events, while many market stress episodes occur without proximate geopolitical triggers. The results suggest that geopolitical risk monitoring should be combined with real-time measures of market fragility to identify when political shocks are likely to become financially systemic.

That version is cleaner and avoids overclaiming.

8. Methodological upgrades I strongly recommend

8.1 Re-estimate the main results excluding COVID as a treatment event

COVID should remain in the sample for market stress detection, but not as a geopolitical conflict shock. Then show:

- HMM detects COVID as systemic stress;
- geopolitical treatment effects survive or change when COVID is not treated as a geopolitical shock;
- Ukraine remains the key geopolitical systemic event.

8.2 Use block randomization inference

Because shocks are clustered and the outcomes are serially persistent, Newey-West alone may not be enough.

Add:

- circular block bootstrap;
- randomization of shock dates within year or month;
- placebo shocks matched on calendar month;
- permutation of shock intensity preserving clustering.

The paper already uses a future-shock placebo, but that is not sufficient.

8.3 Add formal p-values for the state-dependent amplification coefficient

Table 3 reports amplification magnitudes but not p-values or confidence intervals for theta.

You need:

- θ_{hat} ;
- standard error;
- p-value;
- confidence interval;
- total effect at $p = 0.5$ and $p = 0.9$ with standard errors;
- number of identifying shock observations in each region of Pstress.

Without this, a reviewer may say the amplification story is visually suggestive but not formally established.

8.4 Replace or support TCI with a simpler network metric

Because the TCI rolling VAR is vulnerable, add simpler measures:

- rolling average correlation;
- network density above correlation threshold;
- first eigenvalue of the correlation matrix;
- minimum spanning tree length;
- volatility common factor.

If all show the same state-dependent pattern, the paper becomes much more robust.

8.5 Add predictive event classification

Currently the taxonomy is ex post. Add a simple model:

Systemic_i = 1 if event becomes systemic

Predictors:

- pre-event EMFI;
- pre-event TCI;
- pre-event Pstress;
- shock intensity;
- shock breadth;
- event type / global vs country-specific.

This would answer the paper's title question directly:

Which shocks become systemic?

Even a simple logit/probit or classification tree would add substantial value.

9. Internal inconsistencies to fix immediately

These are not deep methodological issues, but they are submission-critical.

1. **Hamas-Israel classification:** abstract says absorbed, Section 7 says systemic. Fix.
2. **EMFI coefficient significance:** Section 6 says $p = 0.174$; conclusion says statistically significant. Fix.
3. **Shock count inconsistency:** the paper mentions 278 shock events, 163 shock-positive trading days, and 154 shock events in the LP/taxonomy. Explain the mapping clearly.
4. **Country sample mismatch:** reconcile the current 19 countries with the companion/BIR sample.
5. **Shock construction description:** remove or align the CAMEO/Goldstein vs GDELT Document conflict-coverage inconsistency.
6. **References:** replace all [?].
7. **Appendices:** currently Appendix headings appear empty in the PDF; that is not submission-ready.
8. **Figure 6 readability:** the network figure is conceptually useful but visually too dense. It needs clearer panels, maybe only pre/post density and top edges, not complete hairball networks.

10. Does this paper now sufficiently differ from BIR?

Yes, much more than before.

The new paper explicitly says the geopolitical shock series is inherited from a companion paper and that shock construction is not the contribution of the current paper. That is exactly the right separation.

The BIR paper remains:

shock construction and validation.

This paper becomes:

systemic market fragility, connectedness, HMM-implied stress, and state-dependent amplification.

That is a defensible two-paper structure.

However, you must keep the shock-construction discussion short and avoid re-claiming it as a methodological contribution. The current paper still includes too much detail and some inconsistent detail. Trim it.

11. My GFJ-oriented recommendation

For Global Finance Journal, I would position the paper as follows:

Best title

When Geopolitical Shocks Become Systemic: Market Fragility and Volatility Connectedness in European Equity Markets

I would remove “Market-Implied Stress” from the title unless the HMM is strengthened. The phrase is interesting but may sound too methodological. A cleaner GFJ title emphasizes the finance mechanism.

Best central claim

Geopolitical shocks do not automatically destabilize markets. They become systemic when they arrive during periods of elevated market fragility, especially when volatility connectedness and tail co-exceedance are already high.

Best contribution

The paper provides a market-based framework for distinguishing absorbed geopolitical shocks from systemic geopolitical stress events.

That is a strong GFJ-style contribution.

12. Final submission advice

Do not submit this exact version yet. It is a strong conceptual draft, but GFJ reviewers would likely raise several serious objections.

Before submission, I would require:

1. Fix country-sample consistency.
2. Remove COVID as a geopolitical treatment event or explicitly reframe it.
3. Resolve the Hamas contradiction.
4. Correct the conclusion’s overstatement of statistical significance.

5. Add formal inference for state-dependent amplification.
6. Strengthen TCI estimation or add simpler network robustness.
7. Make EMFI/HMM real-time or tone down monitoring claims.
8. Add predictive event classification.
9. Replace all citation placeholders.
10. Complete appendices and figure/table documentation.

After these changes, I would view the paper as **potentially suitable for GFJ**. The paper now has a genuinely interesting research question and a distinct contribution relative to BIR. The key is to make the empirical implementation as disciplined as the idea.