

Executive verdict

This version is **substantially stronger** than the previous one and is now much closer to a manuscript that could be credible for **Global Finance Journal**. The paper now has a clear and differentiated research question: **when do geopolitical shocks become systemic financial stress events?** That is exactly the right pivot away from the BIR-style measurement paper.

However, I would still **not submit this exact version yet**. The paper is conceptually much better, but it still has several submission-critical issues. The most serious are not conceptual anymore; they are **internal consistency, empirical overstatement, ex-post versus predictive interpretation, and table/figure coherence**.

My recommendation is:

Potentially suitable for GFJ after one more focused revision, but not ready for submission in the present form.

The idea is now good enough. The manuscript execution still needs tightening.

1. What the paper now does well

The paper has clearly implemented the strategy we discussed. It no longer presents shock construction as the main contribution. Instead, it uses the validated shock series as an external input and builds a new contribution around **market fragility, volatility connectedness, HMM-implied stress states, state-dependent amplification, and systemic-event classification**. The introduction now explicitly states three contributions: a market-based fragility framework, state-dependent geopolitical shock transmission, and a predictive taxonomy of systemic shock outcomes.

This is a major improvement. The new abstract is also much more honest than earlier versions: it says the unconditional local-projection responses are positive but statistically imprecise, and it foregrounds state dependence as the main result.

The paper now has a clear GFJ-style contribution:

Geopolitical shocks are not automatically systemic. They become systemic when they strike already-fragile, highly connected financial markets.

That is a much stronger and more original contribution than the previous claim that geopolitical shocks increase equity volatility.

The paper also now addresses several issues I previously flagged:

Previous issue	Current status
Country sample mismatch	Mostly fixed in the main text: the 19 countries now match the earlier European sample.

Previous issue	Current status
COVID treated as geopolitical shock	Improved: COVID shock days are excluded from the primary treatment series.
Hamas contradiction	Fixed: Hamas-Israel is now consistently classified as systemic with a delayed response.
Overstated unconditional results	Mostly fixed: the paper now says unconditional effects are positive but imprecise.
TCI overparameterization	Partly addressed with wider-window TCI and simpler network metrics.
HMM robustness	Improved with filtered/smoothed comparison and alternative HMM specifications.
Predictive taxonomy	Added and useful.
Empty appendices	Now populated.
Missing references	Mostly fixed.
This is real progress.	

2. Fit with Global Finance Journal

The paper is now **much better aligned with GFJ** than before. It is no longer just a geopolitical-risk event study. It is about:

- European equity-market fragility;
- cross-market volatility connectedness;
- systemic stress states;
- state-dependent shock amplification;
- market-based classification of absorbed versus systemic geopolitical shocks.

That gives the paper a stronger global/international finance angle. The emphasis on **European equity-market integration and systemic fragility** is appropriate for GFJ.

The best GFJ framing is not:

“Geopolitical shocks affect markets.”

It is:

“The financial impact of geopolitical shocks depends on the state of the market system they enter.”

That is interesting, policy-relevant, and finance-oriented.

3. The strongest parts of the current manuscript

3.1 The main research question is excellent

The first paragraph asks: “**when do geopolitical shocks become systemic financial stress events?**” This is a strong and clear question. It gives the paper a real identity.

This is much better than the old paper’s framing, which was too close to “we detect geopolitical shocks and estimate volatility responses.”

3.2 The contribution statement is now much cleaner

The three contributions are now largely correct:

1. a market-based systemic fragility framework;
2. state-dependent shock transmission;
3. a predictive taxonomy of systemic outcomes.

This creates sufficient distance from the BIR paper, provided the shock-series description remains short and consistent.

3.3 The primary empirical finding is now more interesting

The paper no longer tries to force unconditional significance. It correctly states that the unconditional EMFI response is positive but insignificant: $\beta = 0.245$, $p = 0.624$ in the COVID-zeroed primary specification.

That is not a weakness if framed correctly. It supports the paper’s central thesis:

Most geopolitical shocks are absorbed; only some become systemic.

3.4 The state-dependent result is conceptually strong

The smooth-transition specification is the right design. It asks whether the shock effect depends on pre-existing market stress. That is the paper’s core mechanism.

The current Table 3 is much better than the previous version because it now reports standard errors and p-values for the amplification coefficient.

The result is still statistically fragile, but at least the paper is transparent.

3.5 The predictive logit is a very useful addition

The predictive classification section is one of the best additions. It turns the taxonomy from pure ex-post labeling into something closer to an empirical test: can pre-shock market conditions predict which shocks become systemic? The reported leave-one-out AUC of 0.693 and odds ratio of 6.82 for pre-shock EMFI are interesting.

This section should become more central. It answers the title question directly.

4. Submission-critical problems that remain

4.1 The appendices contradict the main text on data construction

This is the most serious remaining issue.

The main text says:

- returns come from **LSEG Workspace**;
- VolStress is the **cross-sectional average of absolute returns**;
- TailVolBreadth counts countries where $|r_{i,t}|$ exceeds a rolling 95th percentile;
- TCI uses a rolling VAR with lag selected by **BIC**.

But Appendix A says:

- prices were downloaded from **Bloomberg** and cross-validated against Refinitiv;
- VolStress is the **cross-sectional median of GARCH(1,1) conditional volatility estimates**;
- TailVolBreadth is based on **conditional volatility**, not absolute returns;
- TCI uses VAR lag selected by **AIC**, not BIC.

This is a major red flag. A reviewer will interpret this as evidence that the manuscript was assembled from inconsistent drafts or that the results are not reproducible.

Fix before submission: Make Appendix A match the main text exactly. If the actual implementation uses absolute returns, remove the GARCH language. If it uses GARCH volatilities, rewrite the main text. The current version cannot be submitted with this inconsistency.

4.2 The shock-series appendix contradicts the main shock description

The main text now correctly says the shock series comes from **GDEL Document 2.0 conflict-related news coverage**, transformed with EVT/GPD and FDR corrections.

But Appendix A.2 says geopolitical shock events are identified using a keyword filter, then aggregated with a **GoldsteinScale-weighted and NumMentions-weighted average**, then scaled by a **min-max transform**.

That is a completely different shock measure. This reintroduces one of the major problems we were trying to eliminate.

Fix before submission: Delete Appendix A.2 or rewrite it to match the BIR/companion shock measure exactly. The paper should not mention GoldsteinScale, NumMentions, min-max scaling, or FDR across event-day outcomes unless those are truly part of the companion shock construction. From the current main text, they are not.

This is not just a cosmetic problem. It affects the credibility of the whole paper.

4.3 The cluster assignments are inconsistent

The main text uses two clusters: Cluster 1 = Western/Northern Europe, Cluster 2 = Central/Eastern and South-Eastern Europe.

Appendix Table 9 uses three clusters and assigns countries differently. For example, Austria is Cluster 1 in the main table but Cluster 3 in Appendix Table 9; Finland is Cluster 1 in the main table but Cluster 2 in Appendix Table 9.

Even if clusters are not central to the paper, this inconsistency signals poor manuscript control.

Fix before submission: Use one cluster scheme or remove clusters entirely. Since clusters are not central to the new paper, I would remove cluster assignments from the main data table unless they are actually used in the results.

4.4 COVID treatment remains conceptually confusing

The paper now excludes 16 COVID-period shock days from the primary LP treatment series because COVID is a systemic financial shock but not a geopolitical conflict shock. That is the correct decision.

But the event taxonomy still classifies **42 of 154 shocks** as systemic, with the systemic category “concentrated in the COVID-19 panic of 2020 and the 2022 Ukraine cluster.”

This creates confusion. If COVID is not a geopolitical conflict shock for the LP, should COVID-related shock days be part of the geopolitical event taxonomy?

There are two clean options:

Option A — preferred

Make the taxonomy consistent with the primary treatment sample:

- primary taxonomy = 138 non-COVID geopolitical shocks;
- COVID shown separately as a market-stress benchmark;
- systemic geopolitical events then focus on Ukraine, Hamas-Israel, and other conflict shocks.

Option B

Keep the 154-event taxonomy but label it clearly:

“Full news-shock taxonomy, including COVID-related shock-news days.”

Then do not call all 154 “geopolitical shocks.”

For GFJ, I recommend Option A. Otherwise a reviewer may argue that the systemic category is partly driven by a non-geopolitical global pandemic.

4.5 The state-dependent inference is still weaker than the wording suggests

The paper says state-dependent projections “reveal” order-of-magnitude amplification. But Table 3 shows:

- EMFI amplification coefficient: $\theta = 1.439$, $p = 0.279$;
- TCI amplification coefficient: $\theta = 2.696$, $p = 0.090$;
- TailVolBreadth: $p = 0.647$;
- Pstress: $p = 0.332$;
- AvgCorr60: negative and insignificant.

So the only conventionally meaningful formal support is for TCI at the 10% level, plus a bootstrap confidence interval excluding zero for TCI. EMFI is economically large but statistically imprecise.

This does not invalidate the paper. But the language should be more precise.

Instead of:

“State-dependent projections reveal that shocks are an order of magnitude more destabilizing.”

Use:

“State-dependent projections provide evidence of economically large amplification, strongest and most precisely estimated for TCI. EMFI amplification is large but imprecisely estimated.”

This distinction matters for GFJ.

4.6 The amplification ratios are unstable because the calm-market denominator is tiny

The paper emphasizes 10.4× and 10.0× amplification. But these ratios are partly mechanical because the calm-market baseline is small: for EMFI, the calm effect is only 0.138.

Ratios can look dramatic when the denominator is close to zero. A reviewer may criticize the “order of magnitude” language as rhetorical.

Recommendation: Report and emphasize **level effects**, not only ratios.

For example:

- calm-market EMFI effect: 0.138;
- near-systemic EMFI effect: 1.433;
- difference: economically meaningful but statistically imprecise;
- TCI near-systemic effect: 2.696 percentage points, more precisely estimated.

Use ratios only as secondary interpretation.

4.7 Some tables and figures still appear to be from older specifications

This is another submission-critical issue.

Table 8 says the primary COVID-zeroed specification has $\beta = 0.245$ and $\theta = 1.44$, while the full-sample version has $\beta = 0.458$ and $\theta = 1.854$.

But Figure 8 still shows the old amplification values: baseline 1.854, EMFI no TCI 1.432, TCI W=60 1.944, TCI W=150 1.554. These appear inconsistent with the new COVID-zeroed primary specification and with Appendix Table 16, where the TCI-window amplification coefficients are much smaller: 0.962, 0.919, 0.761, 0.708, 0.683.

Similarly, Appendix Table 14 reports EMFI $\beta(k=0) = 0.458$, which is the full-sample result, not the primary COVID-zeroed result.

Fix before submission: Regenerate every table and figure from a single consistent pipeline. Label every figure as either:

- primary COVID-zeroed specification;
- full-sample including COVID shock days;
- COVID-excluded sample;
- Ukraine-excluded sample.

Right now, the paper mixes them.

4.8 Figure 4 is unreadable or effectively blank

On page 30, Figure 4 is nearly unreadable in the rendered PDF. The panels are extremely small and the plotted lines are barely visible.

This is a problem because Figure 4 is supposed to show the state-dependent local projection result, which is the paper's core finding.

Fix before submission: Make Figure 4 a full-page figure, preferably with only the three most important outcomes:

1. EMFI;
2. TCI;
3. TailVolBreadth.

Put Pstress and AvgCorr60 in the appendix. The current five-panel version is too dense.

4.9 The predictive logit is useful, but it still has an ex-post construction problem

The paper says the predictive logit uses predictors measured one day before the shock to avoid look-ahead bias. That is partly true. But EMFI is constructed using full-sample standardization and full-sample PCA loadings. The manuscript itself says EMFI is an ex-post outcome variable, not a real-time detector.

Therefore, $EMFI_{t-1}$ in the logit is not fully real-time. It uses future information through the normalization and PCA loadings.

For the LP, this is less problematic because EMFI is an outcome variable. For the **predictive taxonomy**, it matters.

Fix before submission: Estimate one predictive model using a real-time version:

- rolling or expanding standardization;
- PCA loadings fixed from a training window;
- filtered HMM probabilities only;
- leave-one-episode-out or leave-one-year-out validation.

Then report whether AUC remains close to 0.69. If it does, the paper becomes much stronger.

At minimum, call the current logit an **ex-post predictability exercise**, not a real-time predictive model.

4.10 Leave-one-out cross-validation may be too optimistic because events are clustered

The predictive logit uses leave-one-out cross-validation over 154 shock events. But shock events are temporally clustered. Events in the same geopolitical episode are not independent. If the model trains on other days from the same episode, the AUC may be inflated.

Recommendation: Add at least one stricter validation:

- leave-one-year-out;
- leave-one-episode-out;
- block cross-validation by calendar month or quarter;
- cluster-robust bootstrap for AUC.

For GFJ, this would make the predictive section much more credible.

4.11 “Causal” wording should be softened

In Section 8, the manuscript says the future-shock placebo confirms that the positive baseline estimate reflects a “genuine causal response.”

That is too strong. The design is not a clean causal design. The shock series may still be correlated with broader crisis conditions, news cycles, energy shocks, monetary shocks, or unobserved global factors.

Use:

“supports a temporal interpretation”

or:

“reduces concerns that the result is driven by reverse timing or persistent trends.”

Do not say the placebo confirms causality.

5. Remaining GFJ-specific concerns

5.1 The literature review is improved but still too thin

The references are now mostly filled, which is good. But the literature review is still relatively generic. For GFJ, you need a stronger bridge to:

- international equity-market connectedness;
- geopolitical risk and cross-border contagion;
- financial stability and systemic-risk measurement;
- regime switching in global markets;
- European market integration;
- geopolitical risk transmission through energy, FX, and uncertainty channels.

Also, some references appear mismatched. For example, the literature paragraph says recent work using high-frequency intraday data confirms immediate market reactions to geopolitical shocks and cites Caldara and Iacoviello and Rigobon. That is not an obvious or accurate pairing.

Recommendation: Add 8–12 directly relevant recent papers, especially on geopolitical risk and connectedness/spillovers. Also check every reference for fit. GFJ reviewers will expect the paper to be embedded in the international finance literature, not only in methods literature.

5.2 The policy implication is good but should be toned down

The paper concludes that geopolitical risk monitoring should be combined with market-based measures of systemic fragility. That is an excellent policy message.

But because EMFI and HMM are explicitly ex-post, the claim should be framed as:

“These results suggest the value of developing real-time analogues of EMFI and HMM stress probabilities.”

The paper already says real-time counterparts are future work, which is good. But the abstract and introduction should not imply that the current EMFI is already a deployable monitoring tool.

6. Does the paper now sufficiently differ from the BIR paper?

Yes — **much more than before.**

The current manuscript clearly treats the shock series as inherited from Lupu et al. and says shock construction is not repeated. That is the correct separation.

The new paper is about:

- systemic fragility;
- EMFI;
- HMM stress states;
- volatility connectedness;
- state-dependent amplification;
- classification and prediction of systemic shock outcomes.

That is clearly distinct from a BIR-style paper focused on constructing and validating a geopolitical conflict shock measure.

However, the appendix contradiction about shock construction could undermine this distinction. If Appendix A.2 remains, reviewers may think the paper is again trying to define a different shock series. Remove or rewrite it.

7. Recommended revised storyline

I would sharpen the paper into this structure:

Main argument

Geopolitical shocks have weak average effects because most are absorbed. Their systemic consequences are state-dependent: shocks become financially systemic primarily when they arrive in already-fragile, highly connected markets.

Main evidence

1. The unconditional LP effect is positive but statistically weak.
2. TCI shows the strongest state-dependent amplification.
3. EMFI amplification is economically large but imprecise.
4. Pre-shock EMFI predicts systemic classification with moderate out-of-sample power.
5. Most shocks are absorbed; systemic shocks are concentrated in high-fragility windows.

What to avoid

- Do not oversell “10× amplification” as if all outcomes are statistically strong.
- Do not call COVID a geopolitical shock.

- Do not claim causality.
 - Do not imply real-time prediction unless you construct real-time indicators.
-

8. Concrete must-fix list before submission

I would not submit until these are fixed:

1. **Rewrite Appendix A.2** to match the actual companion shock series. Remove GoldsteinScale, NumMentions, min–max scaling if not applicable.
 2. **Reconcile main text and Appendix A.3** on VolStress, TailVolBreadth, and TCI construction.
 3. **Fix data-source inconsistency:** LSEG versus Bloomberg.
 4. **Fix cluster inconsistency:** two clusters in main text, three in appendix.
 5. **Clarify the sample hierarchy:**
 - 278 calendar shocks;
 - 163 trading shock days;
 - 154 post-warm-up shock days;
 - 138 non-COVID geopolitical treatment days;
 - 154-event taxonomy if retained.
 6. **Decide whether taxonomy is 154 full-shock taxonomy or 138 non-COVID geopolitical taxonomy.**
 7. **Regenerate all figures/tables under the same primary specification.**
 8. **Fix Figure 4 readability.**
 9. **Correct Figure 8 values or relabel it as full-sample.**
 10. **Correct Appendix Table 14 and Table 15 labels.**
 11. **Add p-values/CIs for all key state-dependent total effects at $p = 0.9$.**
 12. **Add leave-one-year-out or leave-one-episode-out validation for the logit.**
 13. **Add a real-time or quasi-real-time EMFI robustness for the predictive section.**
 14. **Replace “causal response” with weaker language.**
 15. **Expand and sharpen the literature review for GFJ.**
-

9. Final recommendation

This version is **much closer to a GFJ-ready manuscript**. The research question is now strong, the BIR overlap is much lower, and the paper has a clear finance contribution. The addition of the predictive taxonomy is especially helpful.

But I would still classify it as:

Strong draft, not yet submission-ready.

The main reason is that the manuscript contains serious internal inconsistencies in the appendices and figures. These are exactly the types of issues that can trigger a negative reviewer reaction even when the core idea is good.

After one more revision focused on consistency, table/figure regeneration, and careful wording, I would be comfortable considering this a credible GFJ submission. The paper's best version is not "geopolitical shocks significantly raise fragility on average." The best version is:

Most geopolitical shocks are absorbed. They become systemic only when they enter already-fragile, highly connected markets.

That is the message to build the final submission around.